

Applied Deep Learning - Notes

Dom Hutchinson

October 10, 2020

Contents

1	Machine Learning	2
2	Artificial Neural Networks	2
2.1	Feed-Forward Networks	2
2.1.1	Perceptron	2
2.1.2	Multi-Layer Perceptron	4
3	Training Algorithms	4
3.1	Gradient Descent	5
3.1.1	Auto-Differentiation	5

1 Machine Learning

Definition 1.1 - Deep Representation Learning

Representation Learning is a set of techniques in machine learning where a system can automatically learn representations needed for feature detection from the raw data without the need for hand-designed feature descriptions. *Deep Representation Learning* is then learning to classify using this feature detection.

2 Artificial Neural Networks

Remark 2.1 - Biological Inspiration

In the natural world *Neurons* are the basic working units of the brain. *Neurons* can be split into three main areas

- i). *Dendrites* - Receives inputs from other neurons.
- ii). *Axon* - Carries information.
- iii). *Axon Terminals & Synapses* = Send information to other neurons.

Artificial Neural Networks seek to mimic this structure.

Definition 2.1 - Neuro-Plasticity

Neuro-Plasticity is the ability of a neural system to adapt its structure to accommodate new information (i.e. Learn). This can take several forms including growth & function changes.

2.1 Feed-Forward Networks


Definition 2.2 - Feed-Forward Network

is an artificial neural network where the connections between nodes are uni-directional. Data is provided to the input layer and then an output is returned from the output layer, no layers are visited twice.

2.1.1 Perceptron

Definition 2.3 - Perceptron

A *Perceptron* is an algorithm for supervised learning of a binary classifier. A perceptron defines a hyperplane which acts as a decision boundary which linearly separates the input-state space. These two regions correspond to the two-classes. A perceptron has the following structure.



x_0 is the bias element. It is always set to -1 in the input and the actual value is defined by its weight w_0 .

$\mathbf{x} = (x_0, \dots, x_k)$ is the input. (x_1, \dots, x_k) are the inputs for the item we wish to classify

$\mathbf{w} = (w_0, \dots, w_k)$ is the weights assigned to each input.

Σ is the weighted sum of the bias & inputs. $\Sigma := (\sum_{i=0}^k w_i x_i) = \mathbf{w}^T \mathbf{x}$

g is the *Activation function* which maps from Σ to $\{-1, 1\}$, effectively performing a binary classification. The user has several options for how to define this. (n.b. $g : \mathbb{R} \rightarrow \{-1, 1\}$)

y is the output of the *Activation function*. (i.e. the classification). Typically denoted as $f(\mathbf{x}; \mathbf{w})$

$$y = g\left(\sum_{i=0}^k x_i w_i\right) = g(\mathbf{w}^T \mathbf{x})$$

Remark 2.2 - Limitations of Perceptron

A *Perceptron* can only perform linear binary classification so is not useful when two classes are not linearly separable. See Section 2.2 for how to learn arbitrary decision boundaries.

Proposition 2.1 - Activation Function

There are several choice for the *Activation Function* including:

sign binarily assigns values depend on whether they are positive or negative.

$$\text{sign}(x) := \begin{cases} 1 & x \geq 0 \\ -1 & x < 0 \end{cases} = \frac{x}{|x|}$$

Proposition 2.2 - Perceptron (Supervised) Learning Rule

We need a way for a perceptron to learn when it makes a misclassification. This is done by adjusting the weight vector \mathbf{w} . A simple learning rule is to update the current weights by a certain proportion of the error made.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \Delta \mathbf{w} \quad \text{where} \quad \Delta \mathbf{w} = \begin{cases} \eta \overbrace{f^*(\mathbf{x})}^{\text{ground truth}} \mathbf{x} & \text{if } \overbrace{f^*(\mathbf{x})}^{\text{ground truth}} \neq \overbrace{f(\mathbf{x})}^{\text{prediction}} \\ 0 & \text{otherwise} \end{cases}$$

Here, $\eta \in \mathbb{R}^+$ is know as the *Learning Rate*. Remember that $f^*(\cdot) \in \{1, -1\}$.

Proposition 2.3 - Training Process for a Single-Layer Perceptron

Let $\{(\mathbf{x}_1, f^*(\mathbf{x}_1)), \dots, (\mathbf{x}_N, f^*(\mathbf{x}_N))\}$ be a set of training data. To learn a good set of weights \mathbf{w} do the following process.

- i). Initialise the weight vector $\mathbf{w} = \mathbf{0}$
- ii). Consider next training datum $(\mathbf{x}_i, f^*(\mathbf{x}_i))$.
- iii). Calculate prediction $f(\mathbf{x})$.
- iv). Compare prediction $f(\mathbf{x})$ and ground truth $f^*(\mathbf{x})$.
- v). Update the weight vector $\mathbf{w} = \mathbf{w} + \Delta \mathbf{w}$ where $\Delta \mathbf{w} = \begin{cases} \eta f^*(\mathbf{x}) \mathbf{x} & \text{if } f^*(\mathbf{x}) \neq f(\mathbf{x}) \\ 0 & \text{otherwise} \end{cases}$
- vi). Repeat ii)-v) until the training set is exhausted.


2.1.2 Multi-Layer Perceptron

Remark 2.3 - Learning Arbitrary Decision Boundaries

To learn an arbitrary decision boundary (i.e. anything non-linear) can be done by using a *Multi-Layer Perceptron* with non-linear activation functions.

Definition 2.4 - Multi-Layer Perceptron

A *Multi-Layer Perceptron* has the same general structure as a perceptron but with multiple calculations occurring and multiple output values. Below is a diagram of a MLP of *depth* N (i.e. there are N layers of computation)



Note that each layer can have a different *width* (i.e. number of nodes in the layer). For each consecutive pair of layers $\mathbf{f}^i, \mathbf{f}^j$ (of widths n_i, n_j respectively) there is an associated weight matrix $W \in \mathbb{R}^{n_i \times n_j}$ st $\mathbf{f}^j = W^T \mathbf{f}^i$. The values from the output layer are then passed to an *activation function* to make a classification.

Remark 2.4 - Using MLPs

An MLP with a *single* hidden layer is sufficient to represent any boolean or continuous function, although the layer may be exponentially wider than the input.

An MLP with *two* hidden layers is sufficient to represent any mathematical function.

Proposition 2.4 - MLPs as Computation Graphs

$$\begin{aligned}
 s_i^j &:= (W^j)^T f^{j-1} && \text{weighted sum of the } i^{\text{th}} \text{ node of the } j^{\text{th}} \text{ hidden layer} \\
 \Rightarrow \frac{\partial s_i^j}{\partial w_{ii}^j} &= f_i^{j-1} \\
 f_i^j &:= g_i^j(s_i^j) && i^{\text{th}} \text{ output value of } j^{\text{th}} \text{ hidden layer} \\
 \Rightarrow \frac{\partial f_i^j}{\partial s_i^j} &= \text{depends on def of } g_i^j
 \end{aligned}$$

3 Training Algorithms

Definition 3.1 - Cost/Loss Function, J

A *Cost Function* $J(\cdot; \cdot)$ is a real-valued measure of how inaccurate a classifier is for a given input configuration (test data & weights). Greater values imply the classifier is less accurate. Here are some common cost functions

Expected Loss $J(X; \mathbf{w}) = \mathbb{E}[L(f(\mathbf{x}, \mathbf{w}), f^*(\mathbf{x}))]$

Empirical Risk $J(X; \mathbf{w}) = \frac{1}{|X|} \sum_{\mathbf{x} \in X} L(f(\mathbf{x}, \mathbf{w}), f^*(\mathbf{x}))$

Here $L(x, x^*)$ is a measure of loss (distance) between two values. This is defined by the user on a case by case basis. Popular definitions are: $|x - x^*|$, $(x - x^*)^2$ & $\mathbb{1}\{x = x^*\}$

3.1 Gradient Descent

Definition 3.2 - Gradient Descent

Gradient Descent aims to learn a set of weight values \mathbf{w} which produce a local minimum for a given cost function J . The update rule for gradient descent is

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \underbrace{\eta \cdot \nabla J(X; \mathbf{w}_t)}_{\Delta \mathbf{w}}$$

$\nabla J(X; \mathbf{w}_t)$ is the partial derivative of the cost function wrt to the weights and gives the direction of the greatest descent. We can calculate the i^{th} component of $\Delta \mathbf{w}$ after observing $(\mathbf{x}, f^*(\mathbf{x}))$

$$[\Delta \mathbf{w}]_i = \eta x_i \left(\underbrace{\mathbf{w}_t^T \mathbf{x}}_{f(\mathbf{x}; \mathbf{w}_t)} - f^*(\mathbf{x}) \right)$$

3.1.1 Auto-Differentiation

Proposition 3.1 - Calculating Partial Derivatives

There are three ways to calculate the partial derivatives required for *Gradient Descent*.

- *Symbolic Differentiation* (i.e. algebra). Hard to define to work in all cases.
- *Numerical Differentiation* (i.e. check values in a neighbourhood and approximate the best direction). Easy to implement but low accuracy and high computational cost.
- *Automatic-Differentiation* using feedforward computation graphs. See below

Definition 3.3 - Feedforward Computational Graph


Given a series of equations we can construct a *feedforward computational graph*. *Feedforward computational graphs* have a node for each variable or constant, and then an edge between nodes which are dependent. Once values are defined for all variables at a given depth, values can easily be calculated for variables higher up the tree.

Example 3.1 - Feedforward Computational Graph

Consider the following series of equations

$$\begin{array}{ll} a &= b \times c & b &= d + e \\ c &= e + 2 & d &= 3 + f \\ e &= f \times g \end{array}$$

We can construct the following *Computational Graph*

**Definition 3.4 - Auto-Differentiation using a Feedforward Computational Graph**

Consider two nodes in a computational graph x, y and suppose you want to find the partial derivative $\frac{\partial x}{\partial y}$.

- i). Establish all the paths from y to x in the graph.
- ii). Calculate the partial derivatives of each step of these graphs. (i.e. if there is a path $y \rightarrow a \rightarrow x$ calculate $\frac{\partial a}{\partial y}, \frac{\partial x}{\partial a}$).
- iii). Apply the chain rule along each path (i.e. For $y \rightarrow a \rightarrow x$ calculate $\frac{\partial a}{\partial y} \cdot \frac{\partial x}{\partial a}$).
- iv). Sum these calculations together to get the final result $\frac{\partial x}{\partial y}$.
- v). Substitute variables to make computation easier.

Example 3.2 - Auto-Differentiation using a Feedforward Computational Graph

Consider the graph in **Example 3.1** and wanting to calculate $\frac{\partial f}{\partial a}$.

- i). There are three paths from f to a in the graph: (1) $f \rightarrow d \rightarrow b \rightarrow a$; (2) $f \rightarrow e \rightarrow b \rightarrow a$; and, (3) $f \rightarrow e \rightarrow c \rightarrow a$.
- ii). We need to calculate the following partial derivatives: $\frac{\partial d}{\partial f}, \frac{\partial b}{\partial d}, \frac{\partial a}{\partial b}$ for (1); $\frac{\partial e}{\partial f}, \frac{\partial b}{\partial e}, \frac{\partial a}{\partial b}$ for (2); and, $\frac{\partial e}{\partial f}, \frac{\partial c}{\partial e}, \frac{\partial a}{\partial c}$ for (3).

(1)	(2)	(3)
$\frac{\partial d}{\partial f} = 1$	$\frac{\partial e}{\partial f} = g$	$\frac{\partial e}{\partial f} = g$
$\frac{\partial b}{\partial d} = 1$	$\frac{\partial b}{\partial e} = 1$	$\frac{\partial c}{\partial e} = 1$
$\frac{\partial a}{\partial b} = c$	$\frac{\partial a}{\partial b} = c$	$\frac{\partial a}{\partial c} = b$

- iii). Applying the chain rule to each path gives

$$\begin{aligned}
 (1) \quad \frac{\partial d}{\partial f} \frac{\partial b}{\partial d} \frac{\partial a}{\partial b} &= 1 \cdot 1 \cdot c = c \\
 (2) \quad \frac{\partial e}{\partial f} \frac{\partial b}{\partial e} \frac{\partial a}{\partial b} &= g \cdot 1 \cdot 1 \cdot c = gc \\
 (3) \quad \frac{\partial e}{\partial f} \frac{\partial c}{\partial e} \frac{\partial a}{\partial c} &= g \cdot 1 \cdot b = gb
 \end{aligned}$$

- iv). Summing the terms together we get

$$\frac{\partial a}{\partial f} = c + gc + gb$$

- v). By substitution we get a final expression

$$\frac{\partial a}{\partial f} = 2 + 5g + 2fg + 2fg^2$$

So when $f = 4, g = 2$ we have that $a = 150$ and $\frac{\partial a}{\partial f} = 60$.